

News Topic Classifier Using BERT

Summary:

Objective:

Fine-tune a BERT transformer model to classify news headlines from the AG News dataset into one of four categories: *World*, *Sports*, *Business*, *Sci/Tech*.

Steps Taken:

- **Dataset Loading**
 - Used the AG News dataset via Hugging Face Datasets.
- **Text Preprocessing**
 - Tokenized headlines using bert-base-uncased tokenizer.
 - Applied padding and truncation to standardize input length.
- **Model Fine-Tuning**
 - Fine-tuned the BERT model using Hugging Face's Trainer API.
 - Trained over 3 epochs with a learning rate of $2e-5$.
- **Evaluation Metrics**
 - Evaluated performance using Accuracy and F1-score.
 - Used weighted F1 to handle class distribution fairly.
- **Model Deployment**
 - Deployed a real-time interface using Gradio for live prediction.
 - Users can enter headlines and receive instant topic classification.

What I Learned:

- **Transfer Learning:**

Leveraging pre-trained BERT for a new classification task saves time and improves results.
- **Transformer Tokenization:**

Mastered how to preprocess and prepare text inputs for transformer-based models.
- **Performance Evaluation:**

Gained understanding of why F1-score is critical alongside accuracy in classification tasks.
- **Training with Trainer API:**

Learned how to customize training loops, log metrics, and handle datasets efficiently.

- **Gradio for Deployment:**

Successfully built and deployed a lightweight, interactive UI for live user interaction.

Final Result:

An end-to-end pipeline for news topic classification using BERT, with solid performance and real-world usability from raw text to live predictions.